

Retrieval-Augmented Generation (RAG)

Step 1: Import Required Libraries

What are we performing?

- Importing all the libraries required to build the RAG pipeline.
- Each library provides a specific functionality such as text preprocessing, chunking, embeddings, or similarity search.

Why are we performing it?

- Python cannot use library functions until they are imported.
- These libraries will be used in the upcoming steps of the notebook.

```
In [15]: from langchain_text_splitters import RecursiveCharacterTextSplitter
         from sentence_transformers import SentenceTransformer
         import spacy
         import faiss
```

```
c:\Users\LENOVO\AppData\Local\Programs\Python\Python314\Lib\site-packages\langchain_core\utils\pydantic.py:41: UserWarning:
Core Pydantic V1 functionality isn't compatible with Python
3.14 or greater.
```

```
    from pydantic.v1 import BaseModel as BaseModelV1
```

```
c:\Users\LENOVO\AppData\Local\Programs\Python\Python314\Lib\site-packages\tqdm\auto.py:21: TqdmWarning: IPProgress not found.
Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user\_install.html
    from .autonotebook import tqdm as notebook_tqdm
```

Step 2: Load the Dataset


What are we performing?

- Reading the text file containing Machine Learning statements.
- Storing the complete document in the variable `data` .

Why are we performing it?

- The RAG pipeline requires a knowledge source.
- All preprocessing and embedding generation will be performed on this data.

After this step:



Output

The variable **data** now contains the entire document as one large string.

```
In [2]: data = open(r'machine_learning_2000_sentences.txt').read()
# data
```

\$\$ Text Normalization \$\$

Step 3: Convert Text to Lowercase

What are we performing?

- Converting all characters in the document to lowercase.

Why are we performing it?

- Ensures consistent text representation.
- Prevents words like `Machine` and `machine` from being treated differently during processing.

```
In [3]: data = data.lower()
```

Step 4: Remove Extra Whitespaces

What are we performing?

- Removing multiple consecutive spaces from the document using Regular Expressions (Regex).

Why are we performing it?

- Cleans the text.
- Creates a consistent document for the remaining preprocessing steps.

```
In [4]: import re
data = re.sub(r'\s{2,}',' ',data)
```

Step 5: Remove Statement Numbers

What are we performing?

- Removing labels such as Machine Learning Statement 1: using Regex.

Why are we performing it?

- These labels do not add semantic meaning.
- Prevents unnecessary text from becoming part of the embeddings.

> Syntax:

```
`python
data = re.sub("\d+\.:", "", data)
data
`
```

```
In [8]: # data = re.sub(r'\d+\.','', data)
data = re.sub(r'machine learning statement \d+:','',data)
# data
```

Step 6: Expand Contractions and Abbreviations

What are we performing?

- Expanding contracted words into their full form.
- Example: don't → do not , can't → cannot , ASAP → as soon as possible .

Why are we performing it?

- Creates a standardized text format.
- Improves text quality before further preprocessing.

```
In [9]: import contractions  
data = contractions.fix(data)
```

Step 8: Remove Punctuation

What are we performing?

- Importing Python's built-in punctuation characters.
- Removing all punctuation using Regular Expressions (Regex).

Why are we performing it?

- To identify punctuation symbols that can be removed from the text.
- Keeps only meaningful words.
- Produces cleaner text for NLP preprocessing.

```
In [12]: # import string  
# string.punctuation
```

```
In [11]: data = re.sub(r'^\0-9a-zA-Z\s','', data)
```

Step 9: Spell Correction

What are we performing?

- Importing the TextBlob library for spell correction.
- Applying spell correction to the entire document.

Why are we performing it?

- To correct spelling mistakes in the document before creating embeddings.
- Reduces spelling inconsistencies.
- Improves the quality of embeddings.

```
In [13]: from textblob import TextBlob  
# str(TextBlob(data).correct())
```

Step 10: Load spaCy Model

What are we performing?

- Loading the English language model from spaCy.

Why are we performing it?

- Required for tokenization, stop-word removal, and lemmatization.

```
In [16]: nlp = spacy.load('en_core_web_sm')
```

Step 11: Tokenization and Lemmatization

What are we performing?

- Splitting the text into tokens.
- Converting each word to its root (lemma).
- Removing stop words.

Why are we performing it?

- Reduces unnecessary words.
- Converts similar word forms into a single base word.
- Produces cleaner text for embedding generation.

```
In [18]: tokens = nlp(data)
updated_tokens = [token.lemma_ for token in tokens if not token.is_stop]
# updated_tokens
```

```
In [21]: # if we want to check the word belong to which entity, .ent
# doc = nlp(data)
# for ent in doc.ents:
#     print(ent.text, ent.label_)
```

Step 12: Convert Tokens Back to Text

What are we performing?

- Joining all processed tokens into a single string.
- Removing extra spaces.

Why are we performing it?

- Creates the final cleaned document.
- This cleaned text will be used for chunking.

```
In [22]: data = ' '.join(updated_tokens).strip()
# strip: to remove extra spaces
# we can also use it to remove special characters
```

Step 13: Chunking the Document

What are we performing?

- Preparing to split the document into smaller text chunks.

Why are we performing it?

- LLMs cannot efficiently process very large documents at once.
- Smaller chunks improve retrieval accuracy.

```
In [23]: splitter = RecursiveCharacterTextSplitter(chunk_size = 100,
chunk_overlap = 20)
```

Step 14: Create Text Chunks

What are we performing?

- Splitting the cleaned document into chunks.
- Removing duplicate chunks using `set()` .

Why are we performing it?

- Creates manageable pieces of text.
- Prevents duplicate information from being stored in the vector database.

```
In [27]: chunks = splitter.split_text(data)
# chunks # type of chunks -> list of strings
```

```
In [ ]: # docs = [Document(page_content=text, metadata={"file": "notes.txt"})]
# chunks = splitter.split_documents(data)
# chunks type of chunks -> Document
```

Step 15: Generate Chunk Embeddings

What are we performing?

- Loading the Sentence Transformer model.
- Converting every chunk into a numerical vector (embedding).

Why are we performing it?

- Embeddings capture the semantic meaning of text.
- They allow FAISS to perform similarity search.

What is SentenceTransformer?

SentenceTransformer is a class from the sentence-transformers library.

Its job is to load a pre-trained embedding model.

、
Text
↓
Embedding Model
↓
Numerical Vector
、

```
`python  
SentenceTransformer(  
model_name_or_path="sentence-transformers/all-mpnet-base-v2"  
)  
、
```

- It downloads the model (only the first time) and stores it in the local cache.

What is all-mpnet-base-v2?

It is a pretrained transformer model specifically trained to convert sentences into embeddings.

Example:

"I love machine learning."

↓

[0.23, -0.81, 0.56, ..., 0.14]

、

The output vector has 768 numbers.

```
In [30]: embedding_model = SentenceTransformer(  
          model_name_or_path="sentence-transformers/all-minilm-L6-V2"  
          )
```

```
Loading weights: 100%|██████████| 103/103 [00:00<00:00,  
2209.22it/s]
```

Part 1:

```
`python
embedding_model.encode(chunks)
```

、

- encode() processes each chunk through the transformer model.

Internally:

、

Chunk 1

↓

Transformer

↓

768 numbers

Chunk 2

↓

Transformer

↓

768 numbers

Chunk 3

↓

Transformer

↓

768 numbers

、

Result:

、

```
[
[0.12, -0.56, ..., 0.33],
[0.44, 0.29, ..., -0.87],
[0.18, 0.05, ..., 0.71]
]
```

、

Part 2:

```
`python
.astype("float32")
```

、

- The embeddings produced by `encode()` are often float64 by default.
`0.1234567891234567`
- FAISS requires float32 for efficient similarity search. `0.12345679`
- This reduces memory usage by about half.

In [31]: `chunk_embeddings = embedding_model.encode(chunks).astype('float32')`

Step 16: Find Embedding Dimension

What are we performing?

- Retrieving the size (dimension) of each embedding vector.

Why are we performing it?

- FAISS requires the embedding dimension while creating the vector index.

```

、
chunks = [
    "Machine learning is AI.",
    "Python is easy.",
    "Deep learning uses neural networks.",
    "Data science is fun.",
    "RAG uses embeddings."
]
chunks
↓
embeddings
[
    [0.21, -0.53, 0.84, ..., 0.17], ← Chunk 1
    [0.41, 0.13, 0.66, ..., 0.02], ← Chunk 2
    [0.18, -0.71, 0.44, ..., 0.91], ← Chunk 3
    [0.32, 0.25, 0.19, ..., 0.65], ← Chunk 4
    [0.09, -0.81, 0.57, ..., 0.38] ← Chunk 5
]
、
`python
chunk_embeddings.shape
、
returns (5464, 384)

```

- 5464 → number of chunks (rows)
- 768 → numbers in each embedding (columns)
- For each 5464 chunks, 768 embedding values.

```

In [ ]: # chunk_embeddings.shape (5464, 384)
        dimension = chunk_embeddings.shape[1]
        # chunk_embeddings.shape[1] 384

```

Step 17: Normalize Embeddings

What are we performing?

- Applying L2 normalization to all embedding vectors.

Why are we performing it?

- Ensures all vectors have the same magnitude.
- Improves cosine similarity search in FAISS.

、

Old Vector

[5,12] # length = 13

New Vector

[5/13,
12/13] = [0.385,
0.923]
、

```
In [36]: faiss.normalize_L2(chunk_embeddings)
```

Step 18: Create FAISS Vector Database

What are we performing?

- Creating a FAISS index.
- Storing all chunk embeddings in the index.

Why are we performing it?

- Enables fast similarity search.
- Forms the knowledge base of the RAG system.

Feature	IndexFlatL2	IndexFlatIP
-----	-----	-----
Full name	Index Flat Euclidean Distance	Index Flat Inner Product
Uses	Euclidean (L2) distance	Inner Product (Dot Product)
Goal	Find the smallest distance	Find the largest inner product
Best for	Distance-based search	Cosine similarity (after normalization)
Metric	Euclidean distance	Inner Product
Better score	Smaller distance	Larger score
With normalized vectors	Equivalent to cosine ranking	Equivalent to cosine similarity
Common use in RAG	Less common	Most common
Recommended for Sentence Transformers	Usually no	Yes (after normalization)

1. IndexFlatL2 (Euclidean Distance)

Smaller distance means the vectors are more similar.

2. IndexFlatIP (Inner Product)

```
In [ ]: index_faiss_db = faiss.IndexFlatIP(dimension) # creates an empty FAISS
index
index_faiss_db.add(chunk_embeddings) # copies embedding vectors into
faiss index so they become searchable
```

Step 21: Define Retrieval Function

What are we performing?

- Creating a function to retrieve the most relevant chunks for a query.

Why are we performing it?

- Retrieves relevant context from the vector database before sending it to the LLM.

```

In [ ]: def r_search(query, k=2):
    query_embedding = embedding_model.encode(query).astype('float32')
    query_embedding = query_embedding.reshape(1,-1)
    faiss.normalize_L2(query_embedding)
    distance, index = index_faiss_db.search(query_embedding,k=k)
    R_chunks = [chunks[i] for i in index[0]]
    R_str = ' '.join(R_chunks)
    return R_str

def g_text(r_search):
    import os
    import requests
    from openai import OpenAI

    API_URL = "https://router.huggingface.co/v1/chat/completions"

    headers = {
        'Authorization': f"Bearer {os.environ['HF_TOKEN']}",
    }

    def query(payload):
        response = requests.post(API_URL,
                                headers=headers,
                                json=payload,
                                timeout=90)

        return response.json()

    prompt = f'''
        You're an helpful assistant
        Assigned Task for you:
        Structure my output => {r_search}
    ...

    response = query({
        "messages" : [
            {
                "role": "user",
                "content": f"{prompt}"
            }
        ],
        "model" : "Qwen/Qwen3-0.6B:featherless-ai"
        # USE THESE MODELS
        # DeepSeek-V3
        # Qwen-Instruct
        # Llama Instruct
        # Mistral Instruct

    })
    # return response
    return response["choices"][0]["message"]["content"]

user_prompt = 'Explain Machine Learning?'
user_prompt = re.sub(r'^0-9a-z-A-Z\s',' ', user_prompt)

```

```

r_response = r_search(user_prompt)
g_response = g_text(r_response)
print(g_response)
print()
print(len(r_response))
print(r_response[:500])

```

<think>

Okay, the user wants me to structure my output as a learning statement 1798 workflow, emphasizing overfitting, supporting the use of SVMs. Let me break this down.

First, I need to identify the key points: structure the output, emphasize overfitting, and support SVMs. The original task mentions "learning statement 1798 workflow emphasizing overfitting emphasize support vector machine improve carefully."

I should start by creating a clear structure. Maybe start with the main learning statement, then outline the workflow steps. Emphasize that overfitting is a critical factor and that SVMs are suitable for this. Make sure to mention improving the model's performance through techniques like cross-validation and regularization.

I need to avoid using any tags and keep the language simple. Let me check if I'm following the instructions correctly. No reasoning, no tags. Alright, that should cover it.

</think>

****Learning Statement 1798 Workflow:****

The workflow for learning and applying support vector machines (SVMs) should prioritize addressing overfitting by using techniques such as regularization, cross-validation, and early stopping. SVMs are well-suited for classification tasks due to their ability to handle high-dimensional data and generalize well, but their performance depends heavily on avoiding overfitting through careful model tuning and validation.

107

learning statement 1798 workflow emphasizing overfitting emphasize support vector machine improve carefully

```

In [45]: # Functions to retrieve the most relevant chunks from the FAISS vector
          database.

def r_search(query,k=2): #input query and number of nearest chunks to
    retrieve from FAISS
        # Convert the user's query into an embedding(vector representation)
        query_embedding = embedding_model.encode(query).astype('float32')

        #Reshape the embedding from (768,) to (1,768)
        # FAISS expects a 2D array: (number_of_queries, embedding_dimension)
        query_embedding = query_embedding.reshape(1,-1)

        # Normalize the query vector to unit length
        # Required because IndexFlatIP is being used for cosine similarity
        faiss.normalize_L2(query_embedding)

        # Search the FAISS index for the k most similar vectors
        # distance -> similarity scores
        # index -> indices of the matching vectors
        distance, index = index_faiss_db.search(query_embedding, k = k)

        # Retrieve the original text chunks using the returned indices
        R_chunks = [chunks[i] for i in index[0]]

        # Combine all retrieved chunks into a single string
        R_str = ' '.join(R_chunks)

        # return the retrieved context
        return R_str

# Function to send the retrieved context to an LLM
# and receive a structured response
def g_text(r_search):

    # Used to access environment variables
    import os
    # Used to make HTTP POST request
    import requests

    # Hugging Face Router API endpoint
    API_URL = "https://router.huggingface.co/v1/chat/completions"

    # Authentication header containing your Hugging Face token
    headers = {
        'Authorization' : f'Bearer {os.environ['HF_TOKEN']}',
    }

    # Helper function to send a request to the API
    def query(payload):

        # Send POST request to the Hugging Face API
        response = requests.post(
            API_URL,
            headers=headers,
            json=payload

```

```

    )

    # Convert JSON response into a python dictionary
    return response.json()

# Prompt given to the language model
# Retrieved chunks are inserted here

prompt = f'''
        You're a helpful assistant.
        Assigned Task for you.
        Structure my output => {r_search}
    ...

# Send prompt and model name to API
response = query({
    # Chat messages
    "messages":[
        {
            "role":"user",
            # User prompt sent to the model
            "content": prompt
        }
    ],
    "model" : "deepseek-ai/DeepSeek-R1:novita" # Qwen/Qwen3-
0.6B:featherless-ai
})

# Return Only generated text
return response["choices"][0]["message"]["content"]

# User's question
user_prompt = "Explain Machine Learning."

# Remove all special characters except space
user_prompt = re.sub(r'^0-9a-zA-Z\s',' ',user_prompt)

# Retrieved the most relevant chunks from the FAISS vector database
r_response = r_search(user_prompt)

# send the retrieved chunks to the LLM
# The LLM generates a structured answer using the retrieved context
g_response = g_text(r_response)

# print the final response generated by the LLM
print(g_response)

```

<think>

We are given a task: "Structure my output => learning statement 1798 workflow emphasizing overfitting emphasize support vector machine improve carefully"

Breaking down the task:

1. We are to create a "learning statement" (likely a statement about learning or a lesson learned) with the ID 1798.

2. The workflow should emphasize "overfitting".
3. We must also emphasize "support vector machine" (SVM).
4. The goal is to "improve carefully".

Interpretation:

- We are likely discussing a machine learning workflow that involves SVM and the issue of overfitting.
- The improvement must be done carefully, which might imply using techniques to avoid overfitting while improving the model.

Structure of the output:

- Since it's a structured output, we might need to present it in a specific format. However, the task doesn't specify the exact structure, so we'll create a coherent statement that includes all the required elements.

Key points to cover:

- Mention the workflow (machine learning workflow).
- Emphasize overfitting: explain what it is and why it's a problem.
- Emphasize SVM: explain how SVM works and its strengths.
- Discuss how to improve the SVM model carefully, focusing on avoiding overfitting.

Possible structure for the statement:

1. Introduction: State the context (machine learning workflow) and the model (SVM).
2. Discuss overfitting: Define and explain the risks.
3. Explain SVM: Briefly describe how SVM works and its advantages, but also note that it can be prone to overfitting under certain conditions.
4. Improvement strategy: Describe how we can improve the SVM model while carefully avoiding overfitting (e.g., by tuning hyperparameters, using cross-validation, feature selection, etc.).

Let's draft a statement accordingly:

"Learning Statement 1798: In the machine learning workflow, we emphasize the critical issue of overfitting, which occurs when a model learns the training data too well, including noise, leading to poor generalization on unseen data. We particularly focus on the Support Vector Machine (SVM) algorithm, known for its effectiveness in high-dimensional spaces and robustness against overfitting when properly regularized. To improve SVM performance carefully, we must tune hyperparameters (such as the regularization parameter C and kernel parameters) using cross-validation, select relevant features, and consider ensemble methods to enhance generalization without compromising the model's ability to capture the underlying patterns."

This statement:

- Starts with the learning statement ID.
- Mentions the workflow (machine learning workflow).
- Emphasizes overfitting and SVM.
- Discusses careful improvement strategies.

Note: The task says "emphasize" for both overfitting and SVM, so we have given them prominence. Also, the improvement is done carefully by using techniques that mitigate overfitting.

</think>

Learning Statement 1798: Workflow Emphasis on Overfitting and SVM Improvement

1. Core Focus Areas

- **Overfitting Emphasis**:
 - Explicitly prioritize detecting and mitigating overfitting throughout the workflow.
 - Use techniques like cross-validation, regularization, and pruning to ensure model generalization.
- **Support Vector Machine (SVM) Emphasis**:
 - Leverage SVM's strength in handling high-dimensional data and non-linear boundaries via kernel functions.
 - Highlight SVM's margin-maximization principle to enhance robustness against noise.

**2. Workflow Design

Phase 1: Data Preparation

- **Preprocessing**:
 - Scale features (e.g., standardization) to optimize SVM performance.
 - Apply dimensionality reduction (PCA) to reduce overfitting risks.
- **Train-Test Split**:
 - Stratified sampling to preserve class distribution; reserve 20% for validation.

Phase 2: Model Development

- **SVM Hyperparameter Tuning**:
 - Grid-search `C` (regularization) and `gamma` (kernel influence) to balance bias-variance tradeoff.
 - Validate with **k-fold cross-validation** ($k=10$) to monitor overfitting.
- **Kernel Selection**:
 - Test linear, RBF, and polynomial kernels; log performance to identify overfitting patterns.

Phase 3: Overfitting Mitigation

- **Regularization**:
 - Increase `C` conservatively to avoid overfitting; use L2 penalty as default.
- **Early Stopping**:
 - Track validation loss during training; halt if loss plateaus or increases.
- **Ensemble Integration**:
 - Combine SVM with bagging (e.g., via ensemble methods) to reduce variance.

Phase 4: Evaluation & Improvement

- **Metrics**:
 - Primary: **F1-score** (handles class imbalance).
 - Secondary: **Confusion matrix** analysis.
- **Improvement Strategy**:
 - **Feature Engineering**:

- Iteratively remove low-importance features (using RFE) to simplify the model.
- **Boundary Analysis**:
 - Visualize decision boundaries (using 2D projections) to detect overfitting zones.
- **Adversarial Validation**:
 - Check if validation/test data distributions differ significantly from training data.

3. "Improve Carefully" Protocol

- **Incremental Changes**:
 - Adjust one hyperparameter at a time; document impact on validation metrics.
- **Robustness Checks**:
 - Test on multiple validation sets (e.g., bootstrapped samples) to confirm stability.
- **Fallback Mechanism**:
 - Revert to last generalizable model if new changes degrade validation performance.

4. Key Takeaways

- **SVM-Specific Insight**:
 - > **Kernel choice and C/γ tuning are pivotal: Low C prevents overfitting but risks underfitting; RBF kernels require careful γ calibration to avoid capturing noise.**
- **Overfitting Vigilance**:
 - > **Cross-validation is non-negotiable. If training accuracy $\gg$ validation accuracy, immediately revisit feature engineering or regularization.**

****Final Output Structure**:**

```
```markdown
```

```
[ID: 1798]
```

```
[Emphasis 1: Overfitting] → [Emphasis 2: SVM] → [Improvement:
Careful hyperparameter tuning + validation]
```

```
```
```

This workflow ensures SVM models achieve high generalization while systematically curbing overfitting risks.

```

In [ ]: # Using single function
def rag(query, k=2):

    import os
    import requests

    # Step 1: Convert query to embedding
    query_embedding = embedding_model.encode(query).astype("float32")
    query_embedding = query_embedding.reshape(1, -1)
    faiss.normalize_L2(query_embedding)

    # Step 2: Search FAISS
    distance, index = index_faiss_db.search(query_embedding, k=k)

    # Step 3: Retrieve chunks
    retrieved_chunks = [chunks[i] for i in index[0]]
    retrieved_text = " ".join(retrieved_chunks)

    # Step 4: Create prompt
    prompt = f"""
    You're a helpful assistant.

    Structure the following content:

    {retrieved_text}
    """

    # Step 5: Call LLM
    response = requests.post(
        "https://router.huggingface.co/v1/chat/completions",
        headers={
            "Authorization": f"Bearer {os.environ['HF_TOKEN']}"
        },
        json={
            "messages": [
                {
                    "role": "user",
                    "content": prompt
                }
            ],
            "model": "Qwen/Qwen3-0.6B:featherless-ai"
        },
        timeout=90
    )

    # Step 6: Return generated answer
    return response.json()["choices"][0]["message"]["content"]

user_prompt = 'Explain Machine Learning?'
user_prompt = re.sub(r'^\0-9a-z-A-Z\s', ' ', user_prompt)

r_response = r_search(user_prompt)
g_response = g_text(r_response)
print(g_response)

```

In []:

Exported with [runcell](#) — convert notebooks to HTML or PDF anytime at [runcell.dev](#).